

iml554-classification-assignment03

September 6, 2024

1 Problem statement:

The given dataset is from Dream Club which connects borrowers with investors. We will use lending data from 2007-2010 and build a classifier model to predict whether or not the borrower has paid back their loan in full.

Here are what the columns represent:

The given dataset is from Dream Club which connects borrowers with investors. We will use lending data from 2007-2010 and build a classifier model to predict whether or not the borrower has paid back their loan in full.

Here are what the columns represent1-

credit.policy: 1 if the customer meets the credit underwriting criteria of LendingClub.com, and 0 otherwise.

purpose: The purpose of the loan (takes values “credit_card”, “debt_consolidation”, “educational”, “major_purchase”, “small_business”, and “all_others”).

int.rate: The interest rate of the loan, as a proportion (a rate of 11% would be stored as 0.11). Borrowers judged by LendingClub.com to be more risky are assigned higher interest rates.

installment: The monthly installments owed by the borrower if the loan is funded.

log.annual.inc: The natural log of the self-reported annual income of the borrower.

dti: The debt-to-income ratio of the borrower (amount of debt divided by annual income).

fico: The FICO credit score of the borrower.

days.with.cr.line: The number of days the borrower has had a credit line.

revol.bal: The borrower’s revolving balance (amount unpaid at the end of the credit card billing cycle).

revol.util: The borrower’s revolving line utilization rate (the amount of the credit line used relative to total credit available).

inq.last.6mths: The borrower’s number of inquiries by creditors in the last 6 months.

delinq.2yrs: The number of times the borrower had been 30+ days past due on a payment in the past 2 years.

pub.rec: The borrower’s number of derogatory public records (bankruptcy filings, tax liens, or judgments).

not.fully.paid: whether or not the borrower paid back their loan in full

2 Q-1: Load the dataset and print the metadata in the notebook. 1M

```
[406]: # Importing required packages
import numpy as np
import pandas as pd
import warnings as war
war.filterwarnings("ignore")
```

```
[407]: # Defining dataset csv Path
dataSetPath="C:\\Users\\ASUS\\jupyterworkspace\\Assignment & Mini_
↳Project\\Module_03_Classification\\Assignment\\03_Decision Tree\\loan_data.
↳csv"
# Loading dataSet
dataSetRead=pd.read_csv(dataSetPath)
```

```
[408]: # Displaying first 5 records to confirming data loading
print("*****Displaying below_
↳first 5 records*****")
dataSetRead.head()
```

*****Displaying below first 5 records*****

```
[408]:
```

	credit.policy		purpose	int.rate	installment	log.annual.inc	\
0	1	debt_consolidation	0.1189	829.10	11.350407		
1	1	credit_card	0.1071	228.22	11.082143		
2	1	debt_consolidation	0.1357	366.86	10.373491		
3	1	debt_consolidation	0.1008	162.34	11.350407		
4	1	credit_card	0.1426	102.92	11.299732		

	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	\
0	19.48	737	5639.958333	28854	52.1	0	
1	14.29	707	2760.000000	33623	76.7	0	
2	11.63	682	4710.000000	3511	25.6	1	
3	8.10	712	2699.958333	33667	73.2	1	
4	14.97	667	4066.000000	4740	39.5	0	

	delinq.2yrs	pub.rec	not.fully.paid
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	1	0	0

```
[409]: # Displaying last 5 records to confirming data loading
print("*****Displaying below
↳last 5 records*****")
dataSetRead.tail()
```

```
*****Displaying below last 5
records*****
```

```
[409]: credit.policy      purpose  int.rate  installment  \
9573          0      all_other    0.1461      344.76
9574          0      all_other    0.1253      257.70
9575          0  debt_consolidation  0.1071       97.81
9576          0   home_improvement  0.1600     351.58
9577          0  debt_consolidation  0.1392     853.43

      log.annual.inc    dti    fico  days.with.cr.line  revol.bal  revol.util  \
9573      12.180755  10.39   672      10474.000000      215372      82.1
9574      11.141862   0.21   722      4380.000000        184        1.1
9575      10.596635  13.09   687      3450.041667      10036      82.9
9576      10.819778  19.18   692      1800.000000         0        3.2
9577      11.264464  16.28   732      4740.000000     37879      57.0

      inq.last.6mths  delinq.2yrs  pub.rec  not.fully.paid
9573          2          0          0          1
9574          5          0          0          1
9575          8          0          0          1
9576          5          0          0          1
9577          6          0          0          1
```

```
[410]: # Displaying dimension of dataSet
print("Dimention of Dataset:- {}".format(dataSetRead.shape[0:2]))
print("Total number of rows in Dataset:- {}".format(dataSetRead.shape[0]))
print("Total number of columns in Dataset:- {}".format(dataSetRead.shape[1]))
```

```
Dimention of Dataset:- (9578, 14)
Total number of rows in Dataset:- 9578
Total number of columns in Dataset:- 14
```

```
[411]: # Displaying description & statistical summary of the dataSet
dataSetRead.describe().T
```

```
[411]:          count          mean          std          min  \
credit.policy    9578.0     0.804970     0.396245     0.000000
int.rate         9578.0     0.122640     0.026847     0.060000
installment      9578.0    319.089413    207.071301    15.670000
log.annual.inc   9578.0    10.932117     0.614813     7.547502
dti              9578.0    12.606679     6.883970     0.000000
```

fico	9578.0	710.846314	37.970537	612.000000
days.with.cr.line	9578.0	4560.767197	2496.930377	178.958333
revol.bal	9578.0	16913.963876	33756.189557	0.000000
revol.util	9578.0	46.799236	29.014417	0.000000
inq.last.6mths	9578.0	1.577469	2.200245	0.000000
delinq.2yrs	9578.0	0.163708	0.546215	0.000000
pub.rec	9578.0	0.062122	0.262126	0.000000
not.fully.paid	9578.0	0.160054	0.366676	0.000000

	25%	50%	75%	max
credit.policy	1.000000	1.000000	1.000000	1.000000e+00
int.rate	0.103900	0.122100	0.140700	2.164000e-01
installment	163.770000	268.950000	432.762500	9.401400e+02
log.annual.inc	10.558414	10.928884	11.291293	1.452835e+01
dti	7.212500	12.665000	17.950000	2.996000e+01
fico	682.000000	707.000000	737.000000	8.270000e+02
days.with.cr.line	2820.000000	4139.958333	5730.000000	1.763996e+04
revol.bal	3187.000000	8596.000000	18249.500000	1.207359e+06
revol.util	22.600000	46.300000	70.900000	1.190000e+02
inq.last.6mths	0.000000	1.000000	2.000000	3.300000e+01
delinq.2yrs	0.000000	0.000000	0.000000	1.300000e+01
pub.rec	0.000000	0.000000	0.000000	5.000000e+00
not.fully.paid	0.000000	0.000000	0.000000	1.000000e+00

```
[412]: # Displaying the columns and their respective data types
dataSetRead.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   credit.policy          9578 non-null   int64
1   purpose                9578 non-null   object
2   int.rate               9578 non-null   float64
3   installment            9578 non-null   float64
4   log.annual.inc         9578 non-null   float64
5   dti                   9578 non-null   float64
6   fico                  9578 non-null   int64
7   days.with.cr.line      9578 non-null   float64
8   revol.bal              9578 non-null   int64
9   revol.util             9578 non-null   float64
10  inq.last.6mths         9578 non-null   int64
11  delinq.2yrs            9578 non-null   int64
12  pub.rec                9578 non-null   int64
13  not.fully.paid         9578 non-null   int64
dtypes: float64(6), int64(7), object(1)
```

memory usage: 1.0+ MB

```
[413]: # Handling categorical variable purpose
cat_variable = pd.get_dummies(dataSetRead['purpose']).astype(int)
cat_variable
```

```
[413]:
```

	all_other	credit_card	debt_consolidation	educational	\
0	0	0	1	0	
1	0	1	0	0	
2	0	0	1	0	
3	0	0	1	0	
4	0	1	0	0	
...	
9573	1	0	0	0	
9574	1	0	0	0	
9575	0	0	1	0	
9576	0	0	0	0	
9577	0	0	1	0	

	home_improvement	major_purchase	small_business
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0
...
9573	0	0	0
9574	0	0	0
9575	0	0	0
9576	1	0	0
9577	0	0	0

[9578 rows x 7 columns]

```
[414]: #concatenate the dummies DataFrame with the original one
dataSetRead = pd.concat([dataSetRead,cat_variable],axis=1)
# Drop the original 'purpose' column as it's now redundant
dataSetRead_New = dataSetRead.drop(columns = ['purpose'])
```

```
[415]: # Displaying the columns and their respective data types after handling
↳ categorical variable purpose
dataSetRead_New.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---
```

0	credit.policy	9578	non-null	int64
1	int.rate	9578	non-null	float64
2	installment	9578	non-null	float64
3	log.annual.inc	9578	non-null	float64
4	dti	9578	non-null	float64
5	fico	9578	non-null	int64
6	days.with.cr.line	9578	non-null	float64
7	revol.bal	9578	non-null	int64
8	revol.util	9578	non-null	float64
9	inq.last.6mths	9578	non-null	int64
10	delinq.2yrs	9578	non-null	int64
11	pub.rec	9578	non-null	int64
12	not.fully.paid	9578	non-null	int64
13	all_other	9578	non-null	int32
14	credit_card	9578	non-null	int32
15	debt_consolidation	9578	non-null	int32
16	educational	9578	non-null	int32
17	home_improvement	9578	non-null	int32
18	major_purchase	9578	non-null	int32
19	small_business	9578	non-null	int32

dtypes: float64(6), int32(7), int64(7)
memory usage: 1.2 MB

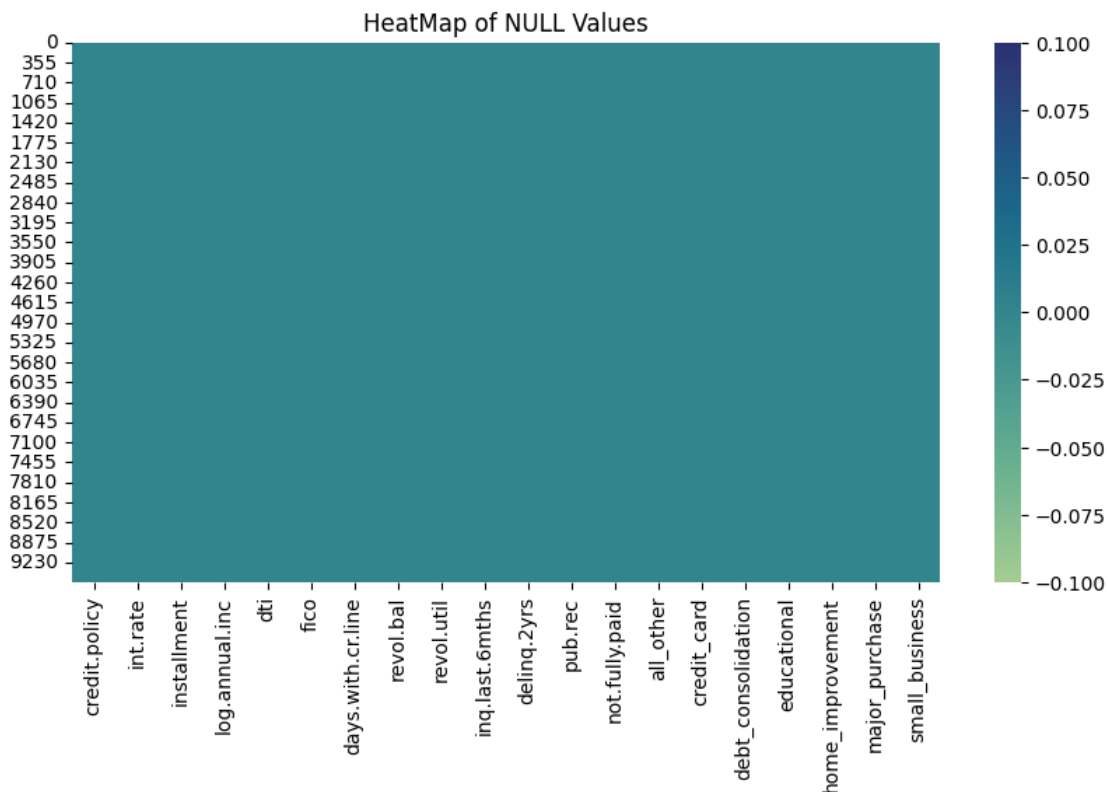
3 Q-2: Print a heatmap to check NULL values. 1M

```
[416]: # Checking total no. of NULL values for attributes specific
missingValue_Count=dataSetRead_New.isnull().sum()
print(missingValue_Count)
```

credit.policy	0
int.rate	0
installment	0
log.annual.inc	0
dti	0
fico	0
days.with.cr.line	0
revol.bal	0
revol.util	0
inq.last.6mths	0
delinq.2yrs	0
pub.rec	0
not.fully.paid	0
all_other	0
credit_card	0
debt_consolidation	0
educational	0
home_improvement	0

```
major_purchase      0
small_business      0
dtype: int64
```

```
[417]: # Import required packages
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(10,5))
sns.heatmap(dataSetRead_New.isnull(),cmap="crest")
plt.title('HeatMap of NULL Values')
plt.show()
```



```
[418]: # Checking duplicate values in dataSet
duplicatevalue_Count=dataSetRead_New.duplicated().sum()
print("Total duplicates values in dataSet:- {}".format(duplicatevalue_Count))
```

Total duplicates values in dataSet:- 0

```
[419]: # Checking percentagewise distiribution for "not.fully.paid" target variable
dataSetRead_New['not.fully.paid'].value_counts(normalize=True).mul(100).round(2)
```

```
[419]: not.fully.paid
0      83.99
1      16.01
Name: proportion, dtype: float64
```

```
[420]: # Cheking countwise distribution for "not.fully.paid" target variable
dataSetRead_New['not.fully.paid'].value_counts()
```

```
[420]: not.fully.paid
0      8045
1      1533
Name: count, dtype: int64
```

Analysis:- Above distribution operation shows that dataset is unbalanced hence SMOTE technique need to be implemented for over sampling

3.1 Split data in to features & target

```
[421]: # Dropping target variable
X=dataSetRead_New.drop('not.fully.paid',axis=1)
# Taking target variable
y=dataSetRead_New['not.fully.paid']
```

```
[422]: # Importing SMOTE module from imblearn library
from imblearn.over_sampling import SMOTE
# Importing Counter module from collections library
from collections import Counter
counter = Counter(y)
print('Before oversampling', counter)
# Oversampling the dataset using SMOTE
smt = SMOTE(random_state = 2)
X_res, y_res = smt.fit_resample(X, y.ravel())
counter = Counter(y_res)
print('After oversampling', counter)
```

Before oversampling Counter({0: 8045, 1: 1533})

After oversampling Counter({0: 8045, 1: 8045})

4 Q-3: Perform stratified splitting of train and test data. 1M

4.1 Splitting the Data into Training and Testing Sets


```
[423]: # Importing train_test_split package
from sklearn.model_selection import train_test_split
# Splitting train & test data
X_train,X_test,y_train,y_test=train_test_split(X_res,y_res,test_size=0.
↪2,random_state=42,stratify=y_res)
```

```
[424]: # Displaying dimension of train & test dataset
print('Shape of X_train = ', X_train.shape)
print('Shape of y_train = ', y_train.shape)
print('Shape of X_test = ', X_test.shape)
print('Shape of y_test = ', y_test.shape)
```

```
Shape of X_train = (12872, 19)
Shape of y_train = (12872,)
Shape of X_test = (3218, 19)
Shape of y_test = (3218,)
```

```
[425]: # Converting a NumPy array to Pandas dataframe
y_test=pd.DataFrame(y_test,columns=['not.fully.paid'])
y_train=pd.DataFrame(y_train,columns=['not.fully.paid'])
```

```
[426]: # Checking distribution for "credit.policy" target variable in y_test
y_test.value_counts()
```

```
[426]: not.fully.paid
0          1609
1          1609
Name: count, dtype: int64
```

```
[427]: # Checking distribution for "credit.policy" target variable in y_train
y_train.value_counts()
```

```
[427]: not.fully.paid
0          6436
1          6436
Name: count, dtype: int64
```

4.2 Data scaling

```
[428]: # importing required package
# MinMax Scaler is used to perform feature scaling
from sklearn.preprocessing import MinMaxScaler
scalling=MinMaxScaler()
scalling.fit(X_train)
```

```
[428]: MinMaxScaler()
```

```
[429]: X_train_scailing=scalling.transform(X_train)
X_test_scailing=scalling.transform(X_test)
```

5 Q-4: Build a classifier model using the Decision Tree algorithm. 2M

```
[430]: # Importing library of DecisionTreeClassifier
from sklearn.tree import DecisionTreeClassifier
# Creating instance of DecisionTreeClassifier
Classifiermodel_DTC_beforePruning = DecisionTreeClassifier()
```

```
[431]: # Training the DecisionTreeClassifier model using training data set
Classifiermodel_DTC_beforePruning.fit(X_train_scailing, y_train)
```

```
[431]: DecisionTreeClassifier()
```

6 Q-5: Print confusion matrix and classification report before and after pruning the Decision tree. (1+1)M

6.1 Printing confusion matrix and classification report before pruning

```
[432]: # Importing Library of classification_report, accuracy_score
from sklearn.metrics import classification_report, accuracy_score
# Importing required packages
from sklearn.metrics import confusion_matrix
```

```
[433]: # Prediction on test data for DecisionTreeClassifier model
y_prediction_DTC_beforePruning=Classifiermodel_DTC_beforePruning.
    ↪predict(X_test_scailing)
```

```
[434]: # Evaluating the accuracy of DecisionTreeClassifier model
accuracy_DTC_beforePruning=accuracy_score(y_test,y_prediction_DTC_beforePruning)
print("DecisionTreeClassifier model accuracy before pruning: {}".
    ↪format(round(accuracy_DTC_beforePruning,2)))
```

DecisionTreeClassifier model accuracy before pruning: 0.8

```
[435]: # Evaluating classification report for DecisionTreeClassifier model
report_DCT_beforePruning=classification_report(y_test,y_prediction_DTC_beforePruning)
print(report_DCT_beforePruning)
```

	precision	recall	f1-score	support
0	0.80	0.79	0.80	1609
1	0.79	0.81	0.80	1609

accuracy			0.80	3218
macro avg	0.80	0.80	0.80	3218
weighted avg	0.80	0.80	0.80	3218

```
[436]: # Evaluating confusion matrix for DecisionTreeClassifier model
conf_matrix_DCT_beforePruning=confusion_matrix(y_test,y_prediction_DTC_beforePruning)
print(conf_matrix_DCT_beforePruning)
```

```
[[1269  340]
 [ 309 1300]]
```

```
[461]: falsePositive = conf_matrix_DCT_beforePruning.sum(axis=0) - np.
        ↪diag(conf_matrix_DCT_beforePruning)
falseNegative = conf_matrix_DCT_beforePruning.sum(axis=1) - np.
        ↪diag(conf_matrix_DCT_beforePruning)
truePositive = np.diag(conf_matrix_DCT_beforePruning)
trueNegative = conf_matrix_DCT_beforePruning.sum() - (falsePositive +
        ↪falseNegative + truePositive)
print('***** DecisionTreeClassifier model before pruning *****')
for i in range(len(truePositive)):
    print(f"Class {i}:")
    print(f"truePositive: {truePositive[i]}, falsePositive: {falsePositive[i]},
        ↪falseNegative: {falseNegative[i]}, trueNegative: {trueNegative[i]}")
    print()
```

```
***** DecisionTreeClassifier model before pruning *****
Class 0:
truePositive: 1269, falsePositive: 309, falseNegative: 340, trueNegative: 1300

Class 1:
truePositive: 1300, falsePositive: 340, falseNegative: 309, trueNegative: 1269
```

6.2 Printing confusion matrix and classification report after pruning

6.3 Pre-Pruning with max_depth

```
[438]: # Train Decision Tree with pre-pruning (using max_depth)
Classifiermodel_DTC_prePruning = DecisionTreeClassifier(random_state=42,
        ↪max_depth=5) # Adjust max_depth
Classifiermodel_DTC_prePruning.fit(X_train_scailing, y_train)
```

```
[438]: DecisionTreeClassifier(max_depth=5, random_state=42)
```

```
[439]: # Prediction on test data for DecisionTreeClassifier model with pre-pruning
y_prediction_DTC_prePruning=Classifiermodel_DTC_prePruning.
↳predict(X_test_scalling)
```

```
[440]: # Evaluating the accuracy of DecisionTreeClassifier model with pre-pruning
accuracy_DTC_prePruning=accuracy_score(y_test,y_prediction_DTC_prePruning)
print("DecisionTreeClassifier model accuracy with pre-pruning: {}".
↳format(round(accuracy_DTC_prePruning,2)))
```

DecisionTreeClassifier model accuracy with pre-pruning: 0.76

```
[441]: # Evaluating classification report for DecisionTreeClassifier model with
↳pre-pruning
report_DCT_prePruning=classification_report(y_test,y_prediction_DTC_prePruning)
print(report_DCT_prePruning)
```

	precision	recall	f1-score	support
0	0.74	0.79	0.76	1609
1	0.77	0.73	0.75	1609
accuracy			0.76	3218
macro avg	0.76	0.76	0.76	3218
weighted avg	0.76	0.76	0.76	3218

```
[442]: # Evaluating confusion matrix for DecisionTreeClassifier model with pre-pruning
conf_matrix_DCT_prePruning=confusion_matrix(y_test,y_prediction_DTC_prePruning)
print(conf_matrix_DCT_prePruning)
```

```
[[1267  342]
 [ 439 1170]]
```

```
[443]: falsePositiveprePruning = conf_matrix_DCT_prePruning.sum(axis=0) - np.
↳diag(conf_matrix_DCT_prePruning)
falseNegativeprePruning = conf_matrix_DCT_prePruning.sum(axis=1) - np.
↳diag(conf_matrix_DCT_prePruning)
truePositiveprePruning = np.diag(conf_matrix_DCT_prePruning)
trueNegativeprePruning = conf_matrix_DCT_prePruning.sum() -
↳(falsePositiveprePruning + falseNegativeprePruning + truePositiveprePruning)
print('***** DecisionTreeClassifier model with prePruning*****')
for i in range(len(truePositiveprePruning)):
    print(f"Class {i}:")
    print(f"truePositive: {truePositiveprePruning[i]}, falsePositive:
↳{falsePositiveprePruning[i]}, falseNegative: {falseNegativeprePruning[i]},
↳trueNegative: {trueNegativeprePruning[i]}")
    print()
```

```
***** DecisionTreeClassifier model with prePruning*****
Class 0:
truePositive: 1267, falsePositive: 439, falseNegative: 342, trueNegative: 1170

Class 1:
truePositive: 1170, falsePositive: 342, falseNegative: 439, trueNegative: 1267
```

7 Cost-complexity pruning (Post-pruning)

```
[444]: Classifiermodel_DTC_postPruning=DecisionTreeClassifier()
path = Classifiermodel_DTC_postPruning.
      ↪cost_complexity_pruning_path(X_train_scailing, y_train)
ccp_alphas, impurities = path.ccp_alphas, path.impurities

[445]: # Train a series of decision trees with different alpha values
pruned_models = []
for ccp_alpha in ccp_alphas:
    pruned_model = DecisionTreeClassifier(criterion="gini", ccp_alpha=ccp_alpha)
    pruned_model.fit(X_train_scailing, y_train)
    pruned_models.append(pruned_model)

# Find the model with the best accuracy on test data
best_accuracy = 0
best_pruned_model = None
for pruned_model in pruned_models:
    accuracy = pruned_model.score(X_test_scailing, y_test)
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_pruned_model = pruned_model
# Model Accuracy after pruning
accuracy_after_pruning = best_pruned_model.score(X_test_scailing, y_test)
print("DecisionTreeClassifier model accuracy post pruning: {}".
      ↪format(round(accuracy_after_pruning,2)))
```

DecisionTreeClassifier model accuracy post pruning: 0.81

```
[466]: # Prediction on test data for DecisionTreeClassifier model post pruning
y_prediction_DTC_postPruning=best_pruned_model.predict(X_test_scailing)

[467]: # Displaying classification report for DecisionTreeClassifier model post pruning
report_DCT_postPruning=classification_report(y_test,y_prediction_DTC_postPruning)
print(report_DCT_postPruning)
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	1609

	1	0.83	0.79	0.81	1609
accuracy				0.81	3218
macro avg	0.81	0.81	0.81	0.81	3218
weighted avg	0.81	0.81	0.81	0.81	3218

```
[468]: # Evaluating confusion matrix for DecisionTreeClassifier model with post pruning
conf_matrix_DCT_postPruning=confusion_matrix(y_test,y_prediction_DTC_postPruning)
print(conf_matrix_DCT_postPruning)
```

```
[[1347  262]
 [ 337 1272]]
```

```
[469]: falsePositivepostPruning = conf_matrix_DCT_postPruning.sum(axis=0) - np.
        ↪diag(conf_matrix_DCT_postPruning)
falseNegativepostPruning = conf_matrix_DCT_postPruning.sum(axis=1) - np.
        ↪diag(conf_matrix_DCT_postPruning)
truePositivepostPruning = np.diag(conf_matrix_DCT_postPruning)
trueNegativepostPruning = conf_matrix_DCT_postPruning.sum() -
        ↪(falsePositivepostPruning + falseNegativepostPruning +
        ↪truePositivepostPruning)
print('***** DecisionTreeClassifier model with Post
        ↪Pruning*****')
for i in range(len(truePositivepostPruning)):
    print(f"Class {i}:")
    print(f"truePositive: {truePositivepostPruning[i]}, falsePositive:
        ↪{falsePositivepostPruning[i]}, falseNegative: {falseNegativepostPruning[i]},
        ↪trueNegative: {trueNegativepostPruning[i]}")
    print()
```

```
***** DecisionTreeClassifier model with Post Pruning*****
Class 0:
truePositive: 1347, falsePositive: 337, falseNegative: 262, trueNegative: 1272

Class 1:
truePositive: 1272, falsePositive: 262, falseNegative: 337, trueNegative: 1347
```

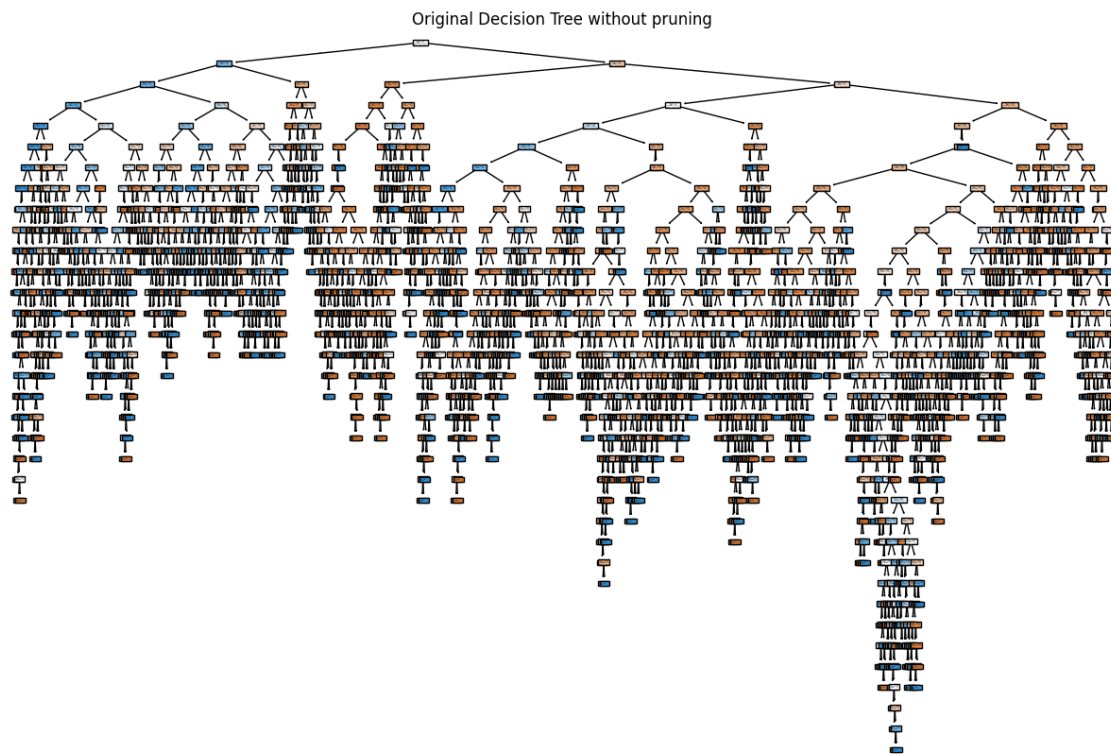
7.0.1 Key Considerations:

Pre-pruning is easier to apply and more straightforward but may result in underfitting if the tree is pruned too early.

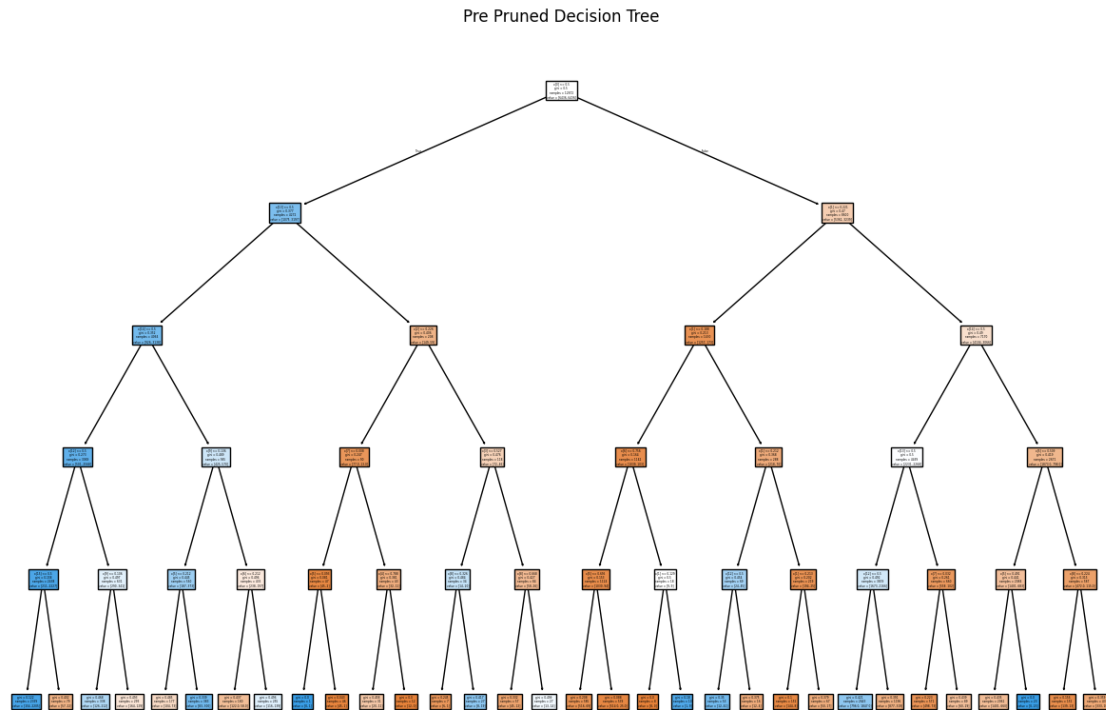
Post-pruning offers a more refined approach by growing the tree fully and then trimming unnecessary branches, but it requires more effort, such as determining the optimal value of `ccp_alpha` through cross-validation.

8 Q-6 Plot the final decision tree model. 1M

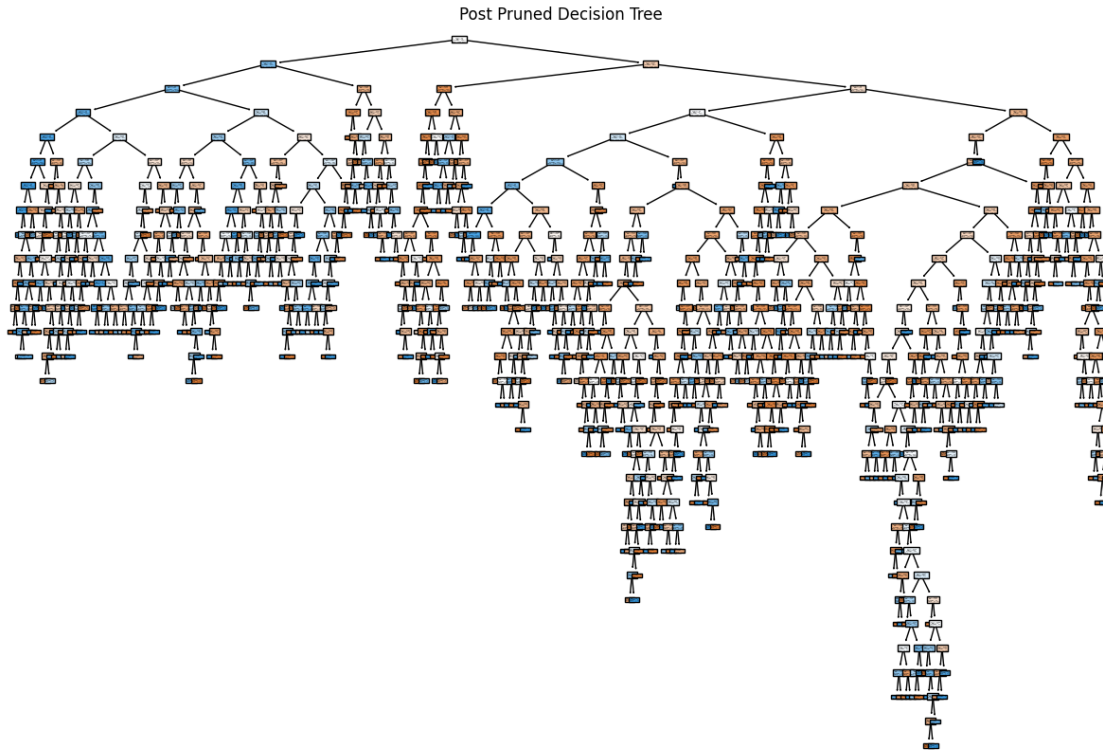
```
[470]: from sklearn.tree import plot_tree
# Plot original tree
plt.figure(figsize=(15, 10))
plot_tree(Classifiermodel_DTC_beforePruning, filled=True)
plt.title("Original Decision Tree without pruning ")
plt.show()
```



```
[471]: # Plot pre pruned tree
plt.figure(figsize=(15, 10))
plot_tree(Classifiermodel_DTC_prePruning, filled=True)
plt.title("Pre Pruned Decision Tree")
plt.show()
```



```
[472]: # Plot post pruned tree
plt.figure(figsize=(15, 10))
plot_tree( best_pruned_model, filled=True)
plt.title("Post Pruned Decision Tree")
plt.show()
```

9 Q-7: Find out the stratified cross-validation accuracy 1M

```
[473]: # Importing required package
from sklearn.model_selection import cross_val_score, StratifiedKFold

[474]: # Initialize the Decision Tree Classifier
Classifier_DST = DecisionTreeClassifier(random_state=42)

# Perform Stratified K-Fold Cross-Validation
# Here, we use 10 folds (cv=10)
stratified_kfold = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)
cross_val_scores = cross_val_score(Classifier_DST, X_res, y_res,
    ↪cv=stratified_kfold, scoring='accuracy')

# Print cross-validation accuracy scores and mean accuracy
print("Cross-Validation Accuracy Scores for each fold:")
print(cross_val_scores)

print("\nMean Stratified Cross-Validation Accuracy:")
print(cross_val_scores.mean())
```

Cross-Validation Accuracy Scores for each fold:

```
[0.81417029 0.80795525 0.80111871 0.79801119 0.81230578 0.82660037  
0.79552517 0.79303915 0.81852082 0.7998757 ]
```

Mean Stratified Cross-Validation Accuracy:
0.8067122436295836